BEYOND MULTIPLE LINEAR REGRESSION · WEEK 2

Treatment effects across centers

Interactions and uncertainty

Dr Douwe Postmus

UMCG · Department of Epidemiology



METHODS · DESIGN · STATISTICS

Week 1 recap: comparing the models

THE REHABILITATION TRIAL

Simulated trial: 120 patients after knee replacement at six centers.

Within each center: 10 patients randomized to each treatment.

Standard care: weekly physiotherapy and home exercises.

Enhanced care: standard care plus a weekly supervised session.

Outcome: six-minute walk distance at 12 weeks.

THE TREATMENT COMPARISON

Model	Estimate (m)	Estimated SE (m)
Treatment only	11.0	5.77
Fixed center effects	11.0	4.34
Random center effects	11.0	4.34

All three models assumed a common treatment effect across centers.

Fixed center parameters and a random center distribution gave the same treatment estimate and estimated standard error in this setting.

Changing treatment allocation within centers

ORIGINAL ALLOCATION

Center	Standard care	Enhanced care
A	10	10
B	10	10
C	10	10
D	10	10
E	10	10
F	10	10
Total	60	60

UNEQUAL ALLOCATION

Center	Standard care	Enhanced care
A	5	15
B	8	12
C	10	10
D	10	10
E	12	8
F	15	5
Total	60	60

Only allocation changes: the underlying treatment effect remains 10 m in every center.

The models now give different treatment estimates

Unequal allocation within centers; estimated enhanced-minus-standard difference in walk distance.

Model	Estimate (m)	Estimated SE (m)
Treatment only	3.87	5.75
Fixed center effects	8.40	4.56
Random center effects	8.13	4.55

Accounting for center changes the estimate substantially. The two center models give slightly different estimates, with almost identical estimated standard errors.

The models use different comparisons

TREATMENT ONLY

The model compares the average outcome of all enhanced-care patients with that of all standard-care patients.

Unequal allocation gives the groups different mixtures of centers. Their difference reflects both treatment and differences in center baselines.

FIXED CENTER EFFECTS

The model allows each center to have its own baseline outcome.

It estimates treatment by combining comparisons between enhanced and standard care within the same center.

These comparisons remove the shared center baseline.

RANDOM CENTER EFFECTS

The model uses comparisons both within and between centers.

Do centers with more enhanced care have higher average outcomes?

Using between-center information assumes baseline center effects are unrelated to treatment allocation.

The random-intercept model also uses between-center information under additional assumptions, so its treatment estimate can differ.

Fixed versus random center effects

Design	Treatment estimates	Estimated standard errors
Equal center sizes; 50:50 allocation within each center	Identical	Identical
Unequal center sizes; 50:50 allocation within each center	Identical	Can differ
Treatment proportions differ across centers	Can differ	Can differ

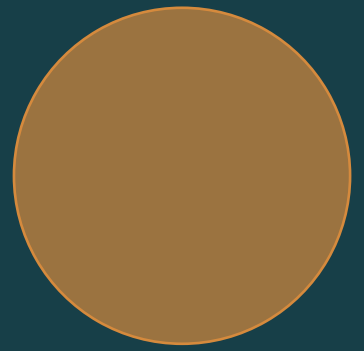
With 50:50 allocation within every center, all three models, including the treatment-only model, give the same treatment estimate, even when center sizes differ.

So far, we have assumed the same treatment effect in every center. What changes when the treatment effect varies across centers?

PART 1

Treatment-by-center interaction

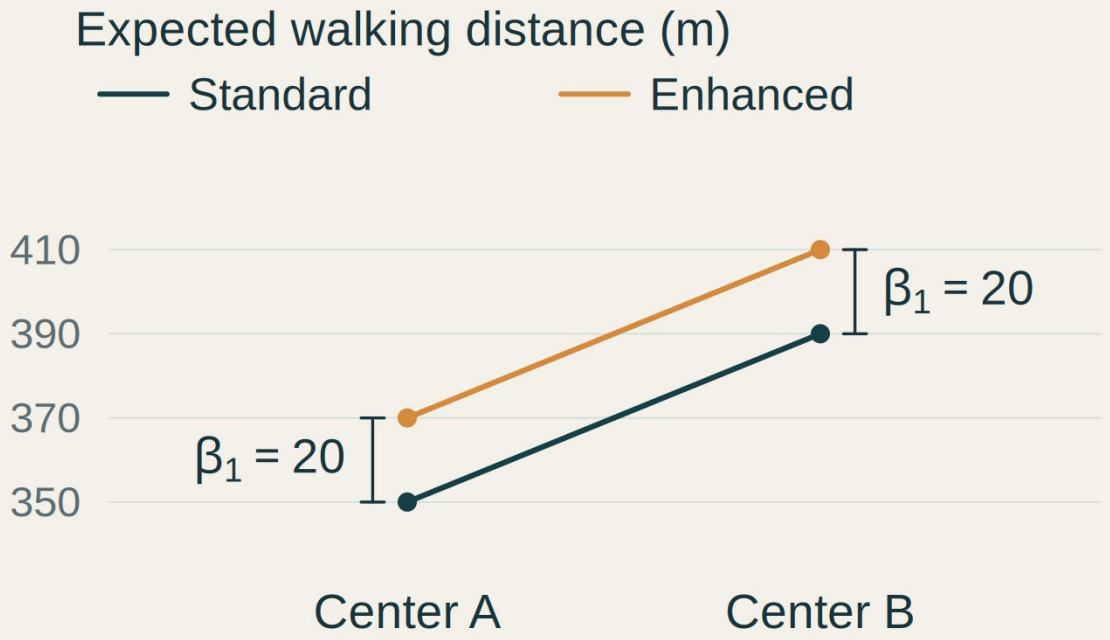




A common treatment effect in both centers

Hypothetical example: assume the four true group means are known.

Center	Standard	Enhanced	Row average
A	$\mu_{0A} = 350$	$\mu_{1A} = 370$	**360**
B	$\mu_{0B} = 390$	$\mu_{1B} = 410$	**400**
Column average	$\beta_0 = 370$	390	**380**



μ_{tj}: expected outcome for treatment t in center j

$$Y_{ij} = \underbrace{\beta_0 + c_j}_{\text{Standard-care mean in center } j} + \underbrace{\beta_1 T_{ij}}_{\text{Treatment contribution}} + \underbrace{\varepsilon_{ij}}_{\text{Individual deviation}}$$

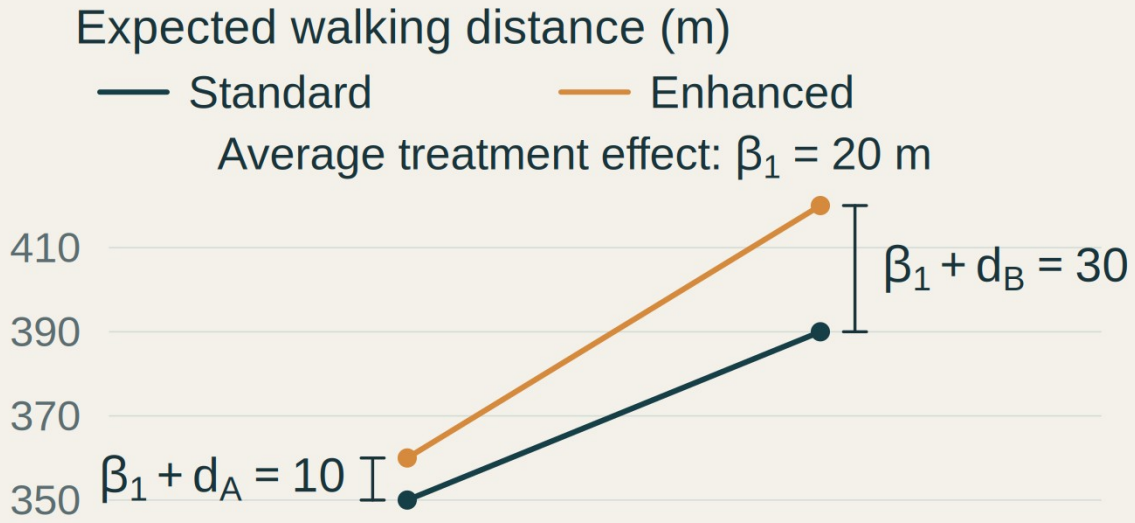
T_{ij} = 0 for standard care and 1 for enhanced care. i = patient; j = center.

Individual patient outcomes vary around their group mean through ε_{ij}.

Different treatment effects across centers

Hypothetical example: assume the four true group means are known.

Center	Standard	Enhanced	Row average
A	$\mu_{0A} = 350$	$\mu_{1A} = 360$	**355**
B	$\mu_{0B} = 390$	$\mu_{1B} = 420$	**405**
Column average	$\beta_0 = 370$	390	**380**



μ_{tj}: expected outcome for treatment t in center j

$$Y_{ij} = \underbrace{\beta_0 + c_j}_{\text{Standard-care mean in center } j} + \underbrace{(\beta_1 + d_j)T_{ij}}_{\text{Treatment contribution}} + \underbrace{\varepsilon_{ij}}_{\text{Individual deviation}}$$

T_{ij} = 0 for standard care and 1 for enhanced care. i = patient; j = center.

Individual patient outcomes vary around their group mean through ε_{ij}.

How should we represent differences between centers?

Beyond two centers, we still allow differences in baseline outcome and treatment effect.

FIXED CENTER EFFECTS

Give each included center its own baseline deviation and treatment-effect deviation.

Express the average treatment effect across the included centers.

RANDOM CENTER EFFECTS

Describe baseline deviations and treatment-effect deviations through a population distribution.

Express the average treatment effect in the population represented by those centers.

We start with fixed center effects and ask what we can learn about treatment in the included centers.

From known means to patient data

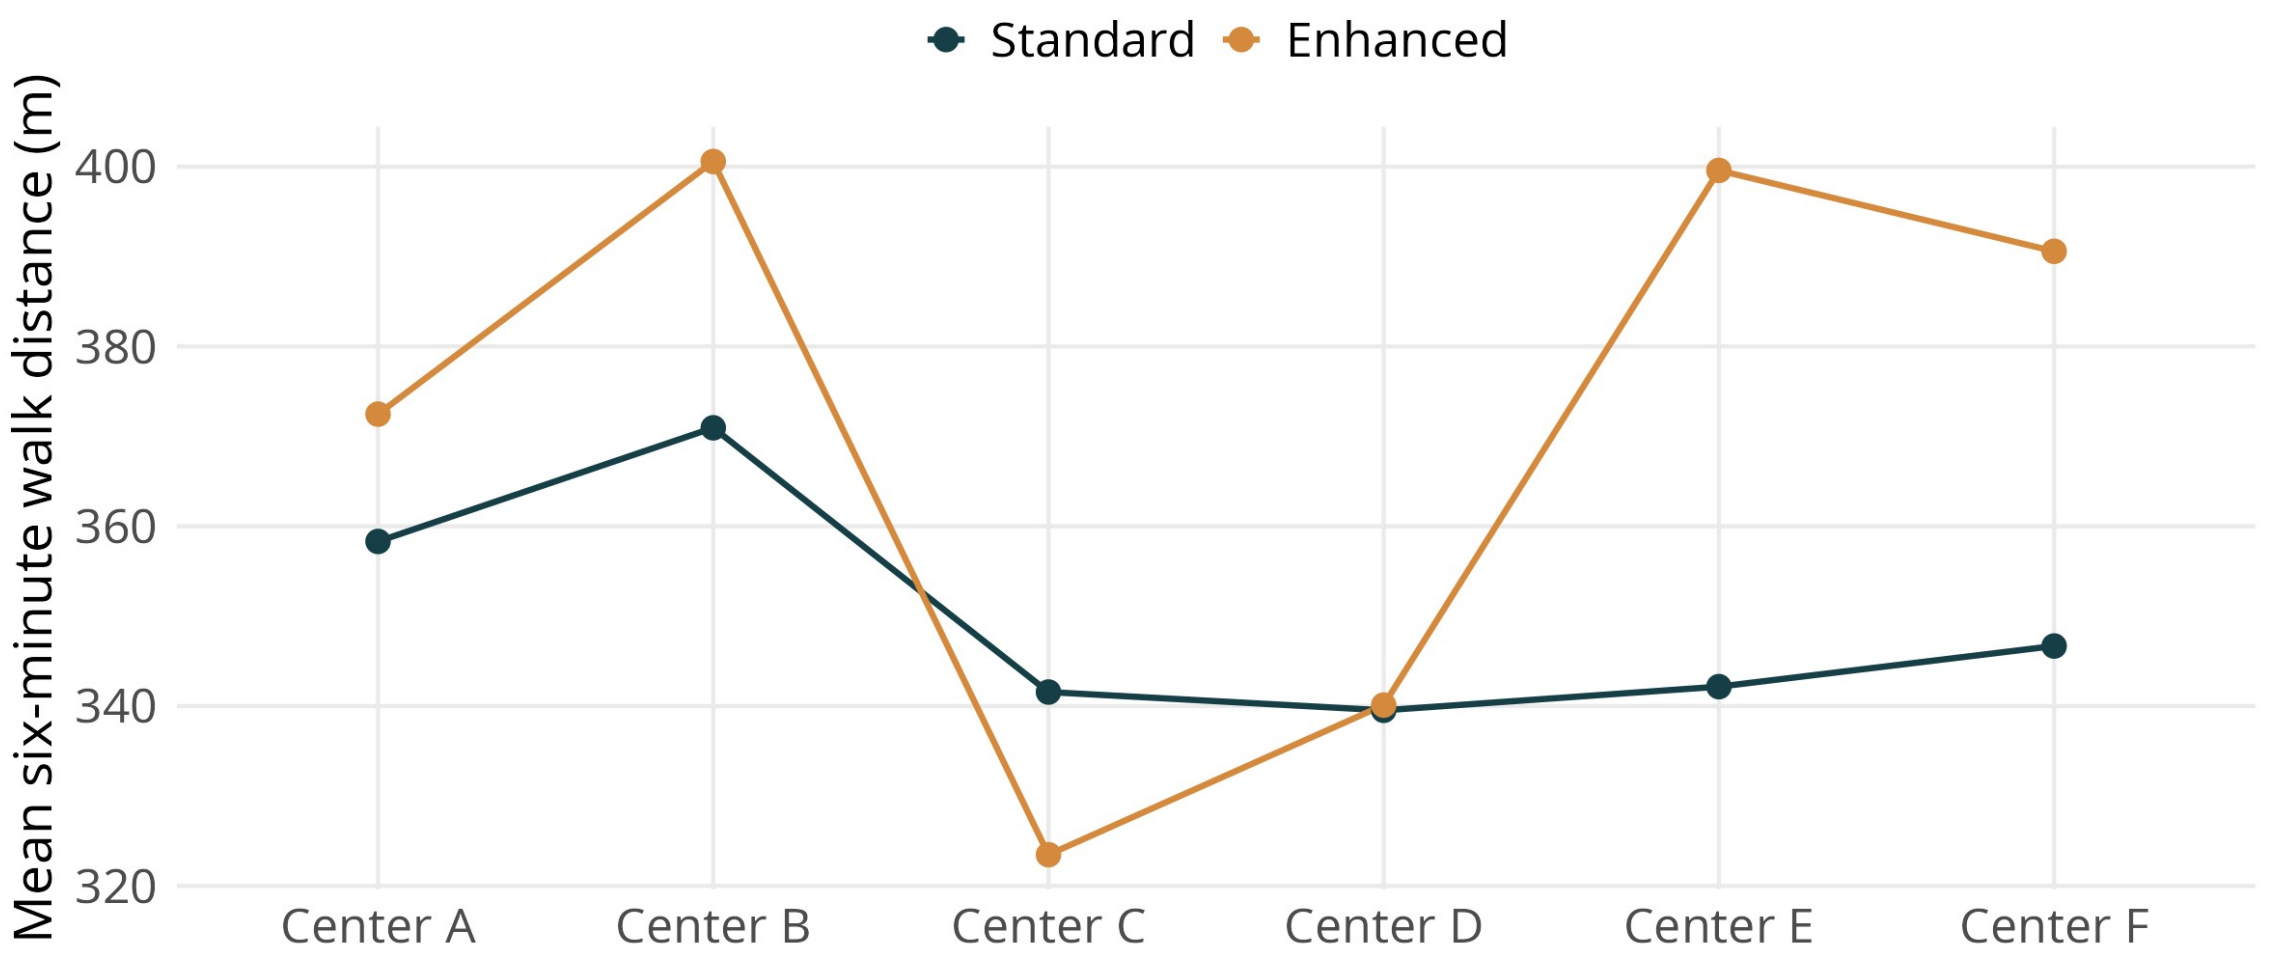
THE REHABILITATION TRIAL

- New simulated trial: 120 patients after knee replacement at six centers.
- Within each center: 10 patients randomized to standard care and 10 to enhanced care.
- Outcome: six-minute walk distance at 12 weeks.
- Fit the fixed-effects interaction model to estimate treatment effects in these centers.

The group means are now unknown and must be estimated from patient outcomes.



Explore outcomes by treatment and center



A fixed-effects model for treatment and center

$$Y_{ij} = \beta_0 + c_j + (\beta_1 + d_j)T_{ij} + \varepsilon_{ij}$$

T_{ij} = 0 for standard care and 1 for enhanced care. Center uses effects coding.

BASELINE OUTCOMES

β_0: average standard-care outcome across the included centers.

c_j: center j's deviation from that baseline average.

$$\sum_j c_j = 0$$

TREATMENT EFFECTS

β_1: average treatment effect across the included centers.

d_j: center j's deviation from the average treatment effect.

$$\sum_j d_j = 0$$

Each center receives equal weight in the averages. Center deviations are fixed, unknown parameters.

Fixed-effects interaction model: results

Term	Estimate (m)	Estimated SE (m)	95% CI (m)	p-value
Intercept (β_0)	349.9	3.31	343.3 to 356.4	< .001
Enhanced care (β_1)	21.3	4.68	12.0 to 30.6	< .001
Center A	8.4	7.41	-6.2 to 23.1	0.257
Center B	21.1	7.41	6.4 to 35.8	0.005
Center C	-8.3	7.41	-23.0 to 6.4	0.265
Center D	-10.3	7.41	-25.0 to 4.3	0.166
Center E	-7.7	7.41	-22.4 to 7.0	0.301
Enhanced × Center A	-7.1	10.48	-27.9 to 13.7	0.499
Enhanced × Center B	8.4	10.48	-12.4 to 29.1	0.427
Enhanced × Center C	-39.4	10.48	-60.1 to -18.6	< .001
Enhanced × Center D	-20.7	10.48	-41.4 to 0.1	0.051
Enhanced × Center E	36.2	10.48	15.4 to 56.9	< .001

Effects coding: Center F follows from the sum-to-zero constraints.

Testing interaction by comparing two models

Fit the full model → fit the reduced model → compare their fit.

Model	Residual sum of squares (RSS)	Residual df
Reduced: common treatment effect	90,694	113
Full: center-specific treatment effects	71,112	108

RSS: residual sum of squares. df: residual degrees of freedom.

$$F = \frac{(\mathrm{RSS}_{\mathrm{reduced}} - \mathrm{RSS}_{\mathrm{full}})/(\mathrm{df}_{\mathrm{reduced}} - \mathrm{df}_{\mathrm{full}})}{\mathrm{RSS}_{\mathrm{full}}/\mathrm{df}_{\mathrm{full}}} = \frac{(90{,}694 - 71{,}112)/(113 - 108)}{71{,}112/108} \approx 5.95$$

Numerator: improvement in fit per additional parameter.

Denominator: estimated individual-error variance from the full model.

F(5, 108) = 5.95, p < .001: strong evidence that treatment effects differ across centers.

Which treatment effect do we want to estimate?

SIMPLE TREATMENT EFFECTS

How do treatments compare within a particular center?

Compare the enhanced and standard means within each center.

$$\text{Effect in center } j: \quad \beta_1 + d_j$$

One treatment effect per center.

MARGINAL TREATMENT EFFECT

How do treatments compare on average across centers?

Average each treatment's mean equally across centers, then compare the averages.

$$\text{Average effect:} \quad \beta_1$$

With effects coding, the treatment deviations average to zero.

Estimated marginal means (EMMs) are model-based estimates of group means, averaged over other factors where appropriate. Comparing these means gives the treatment effects.

Estimating simple treatment effects with estimated marginal means

CODING AND EFFECTS

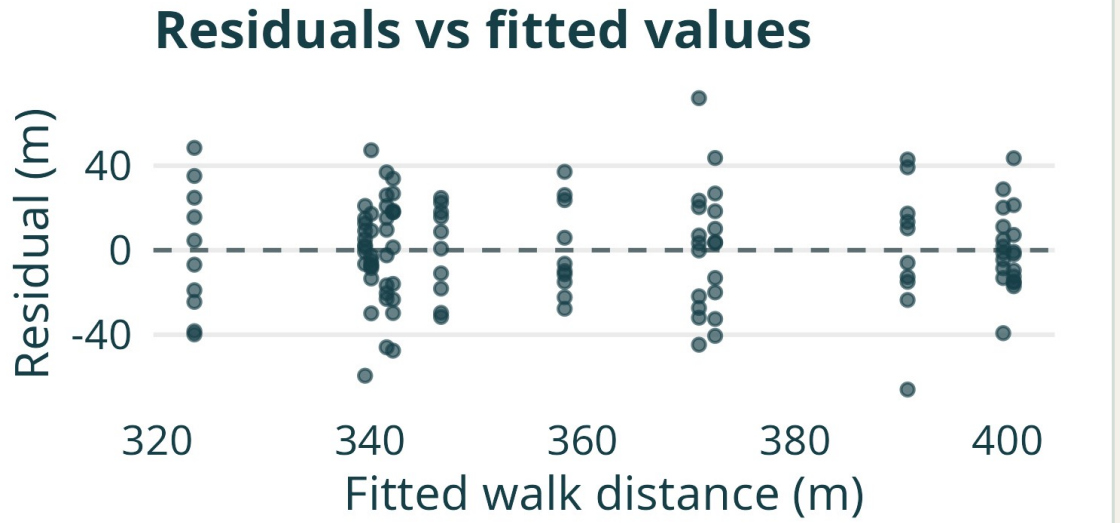
- Coding determines which effects appear directly as coefficients.
- Here, a center's effect combines the average effect with its treatment deviation.
- Predict both treatment means within each center, then compare them.

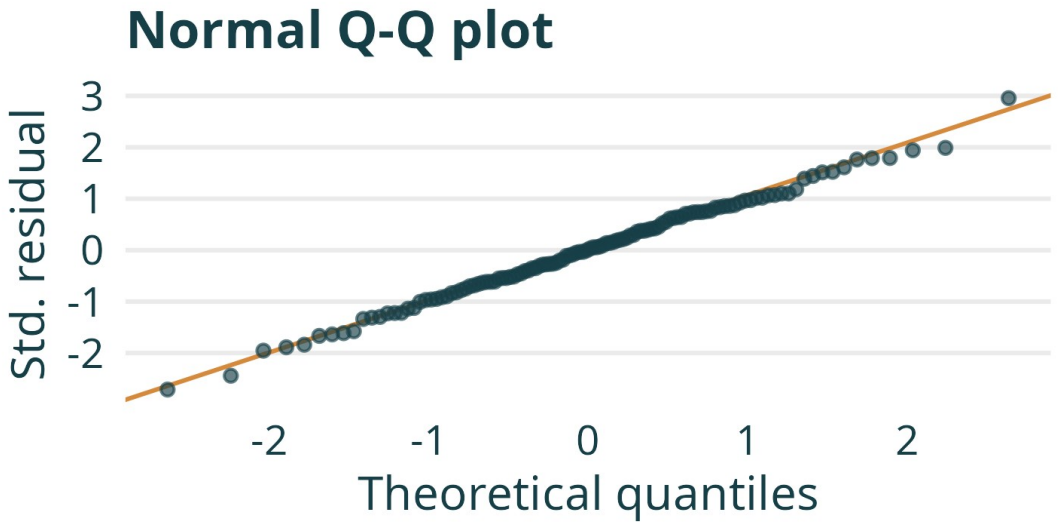
	Predicted means		Simple treatment effect		
Center	Standard mean	Enhanced mean	Difference	95% CI	p-value (Bonf.)
A	358.3	372.5	14.2	-8.6 to 36.9	1.000
B	370.9	400.6	29.6	6.9 to 52.4	0.067
C	341.6	323.5	-18.1	-40.8 to 4.7	0.707
D	339.5	340.1	0.6	-22.2 to 23.3	1.000
E	342.2	399.6	57.4	34.7 to 80.2	< .001
F	346.7	390.6	43.9	21.2 to 66.6	0.001

Means, differences and CIs in meters. Difference: enhanced minus standard. Pointwise 95% CIs; p-values adjusted across six comparisons (Bonferroni).

Equivalent coding schemes give different coefficients, but the same fitted means and treatment comparisons.

Model diagnostics: fixed-effects interaction





WHAT WE USE THIS PLOT FOR

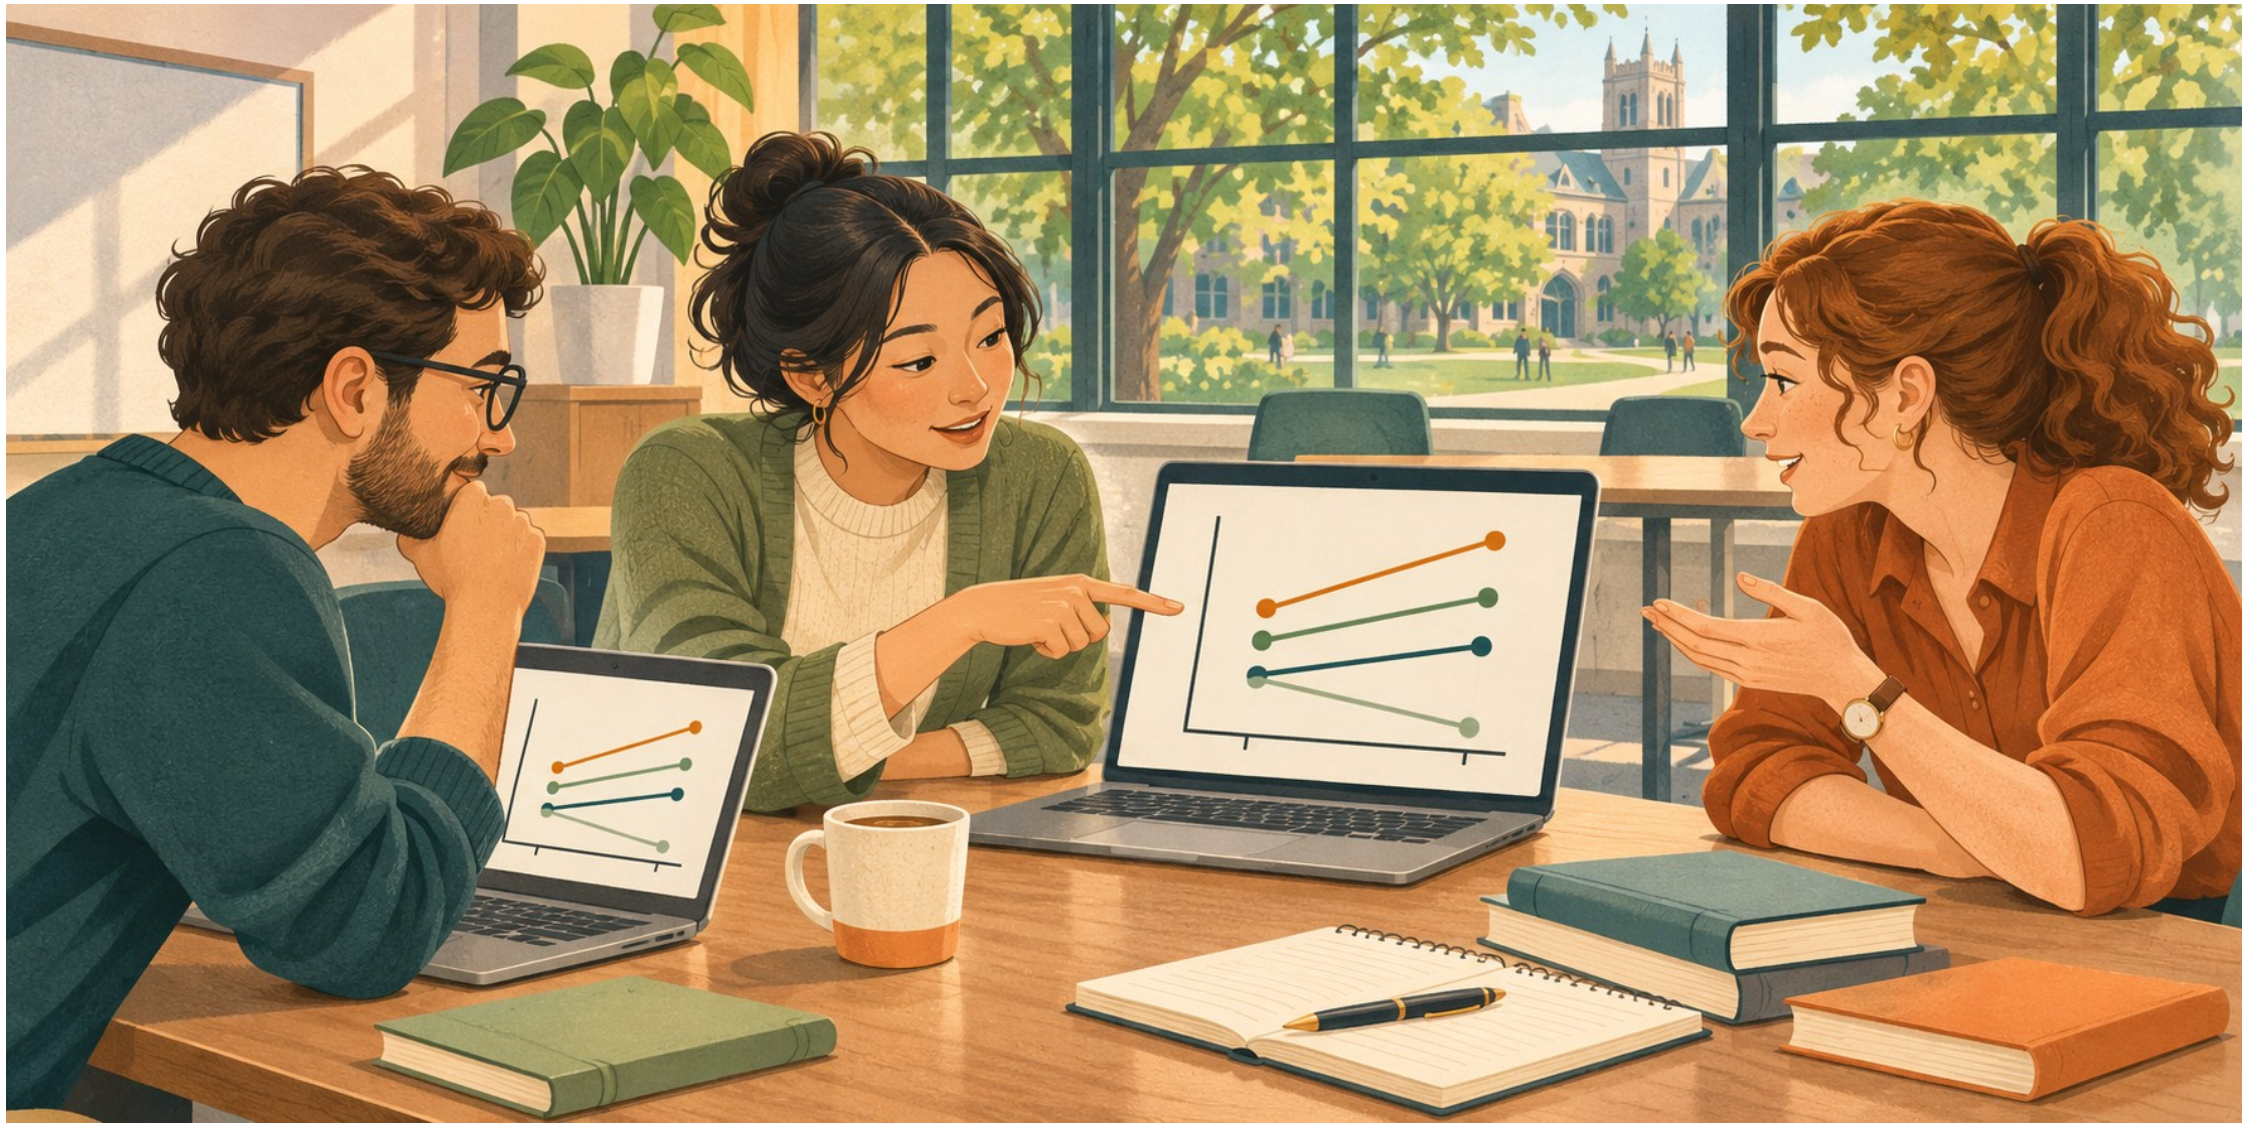
Assess whether residual spread is similar across fitted walk distances.

WHAT WE USE THIS PLOT FOR

Assess whether the residual distribution is approximately Normal.

The plots show no clear change in residual spread and only small departures from the Normal reference line.

Time for Lab 3



Lab 3: what if we leave out the interaction?

Average effect across the five laboratories
Experimental minus standard, in percentage points of wound closure

Model	Estimate	Estimated SE	95% CI	p-value
Additive	9.28	1.81	5.69 to 12.88	< .001
With interaction	9.28	1.62	6.07 to 12.49	< .001

Same estimate: 10 plates per treatment in every laboratory make both models estimate the average of the five within-laboratory treatment differences.

Larger estimated SE: without interaction, differences in treatment effects remain in the residuals, increasing the estimated residual variation.

In this balanced study, omitting interaction leaves the average treatment estimate unchanged but increases its estimated uncertainty.

What if laboratory sizes differ?

WHAT THE MODELS ESTIMATE

- With our effects coding, the interaction model estimates the equally weighted average treatment effect, regardless of laboratory sizes.
- If laboratories have equal sizes, the additive model estimates this same average.
- If laboratory sizes differ, the additive model gives larger laboratories more weight.

HYPOTHETICAL EXAMPLE

Effects in percentage points

Laboratory	Plates	Treatment effect
Small	20	5
Large	80	15

Interaction model: 10 points

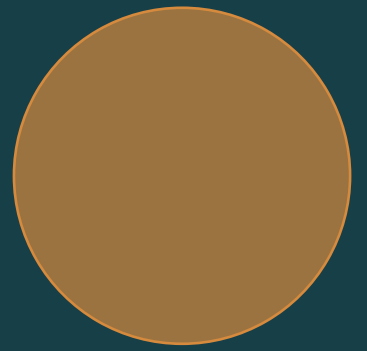
Additive model: 13 points

When treatment effects differ across laboratories, omitting interaction can change both the interpretation of the treatment coefficient and its estimated uncertainty.

PART 2

From included centers to a population





From the included centers to a population

THESE SIX CENTERS

Imagine repeating the study with new patients in the same six centers.

The centers' true treatment effects stay the same.

We estimate the average effect in these six centers.

A POPULATION OF CENTERS

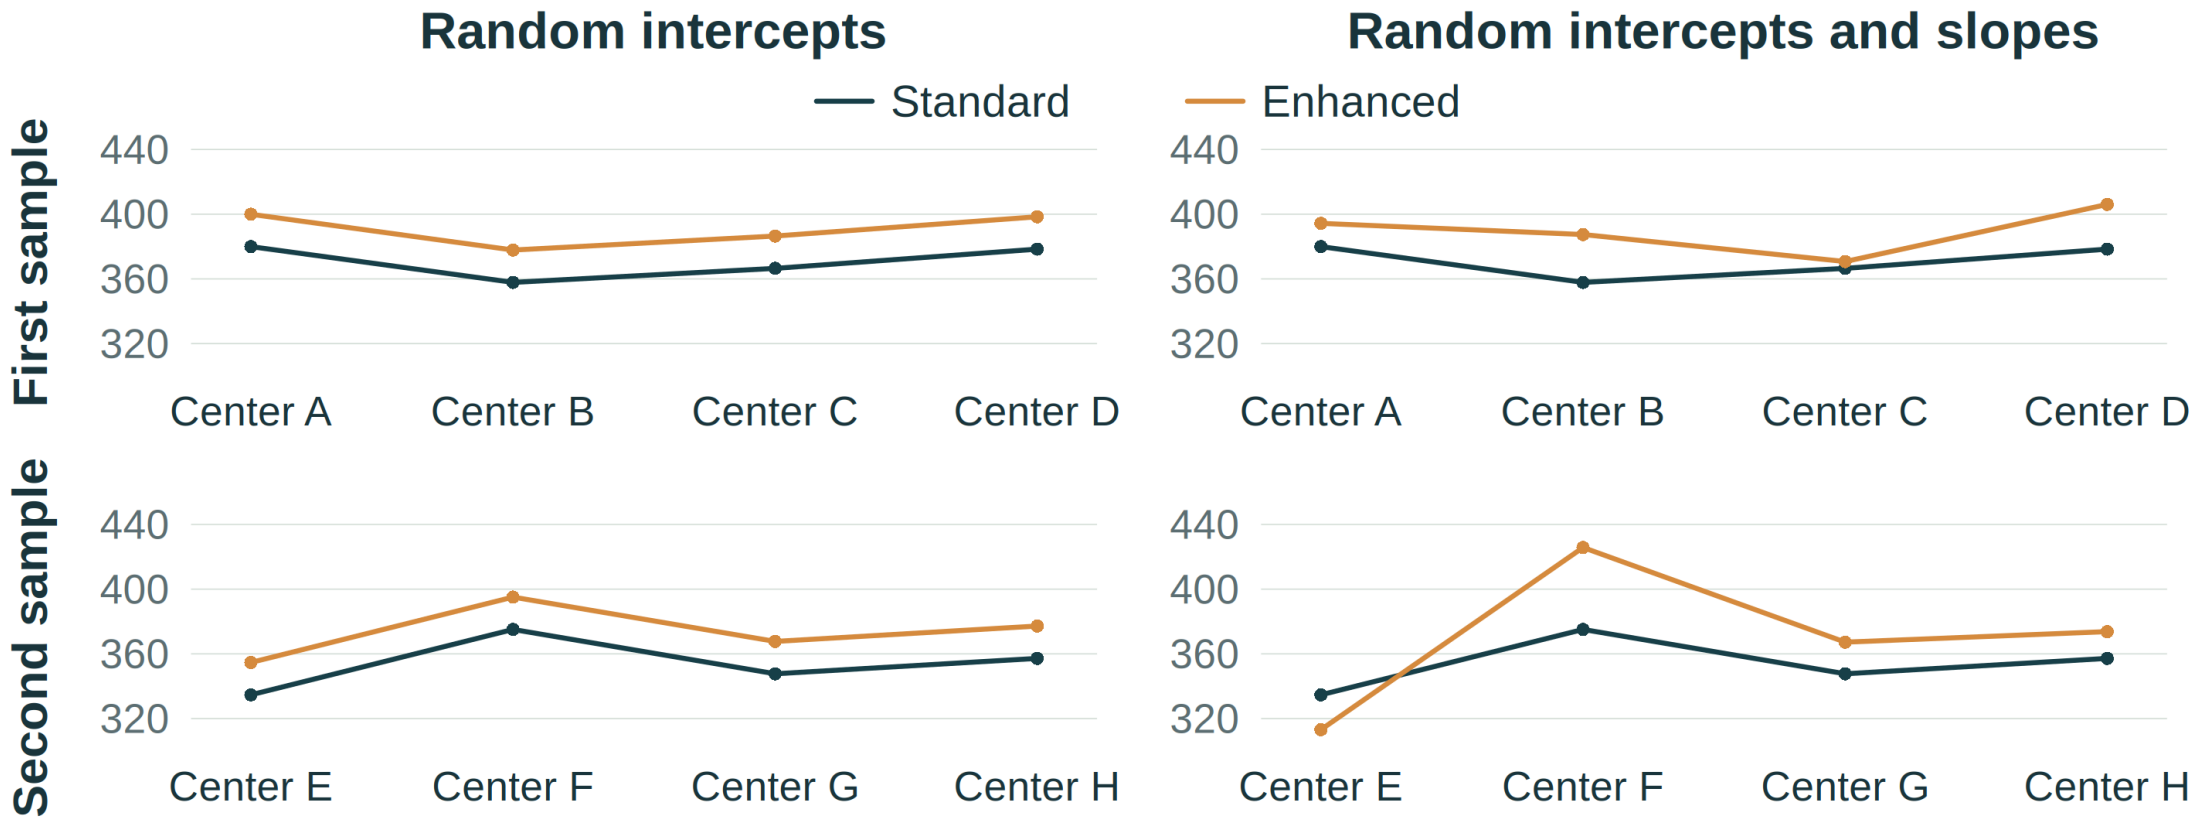
Imagine repeating the study with new centers from the same population.

The selected centers can have different true treatment effects.

We estimate the average effect across the population of centers.

When we generalize to a population of centers, uncertainty also depends on which centers enter the study.

What changes when we sample new centers?



Hypothetical true mean walking distances (m), without patient sampling noise.

The population model stays the same, but each sample contains centers with different realized effects.

The center effects follow a population distribution

$$Y_{ij} = \beta_0 + c_j + (\beta_1 + d_j)T_{ij} + \varepsilon_{ij}$$

β_0 and β_1 are population averages. T_{ij} = 0 (standard) or 1 (enhanced).

DISTRIBUTIONAL ASSUMPTIONS

$$\begin{pmatrix} c_j \\ d_j \end{pmatrix} \sim N\left[\begin{pmatrix} 0 \\ 0 \end{pmatrix}, \begin{pmatrix} \sigma_c^2 & \sigma_{cd} \\ \sigma_{cd} & \sigma_d^2 \end{pmatrix}\right]$$

$$\varepsilon_{ij} \sim N(0, \sigma^2)$$

Effect pairs are independent across centers. Errors are mutually independent and independent of the center effects.

VARIANCES AND COVARIANCE

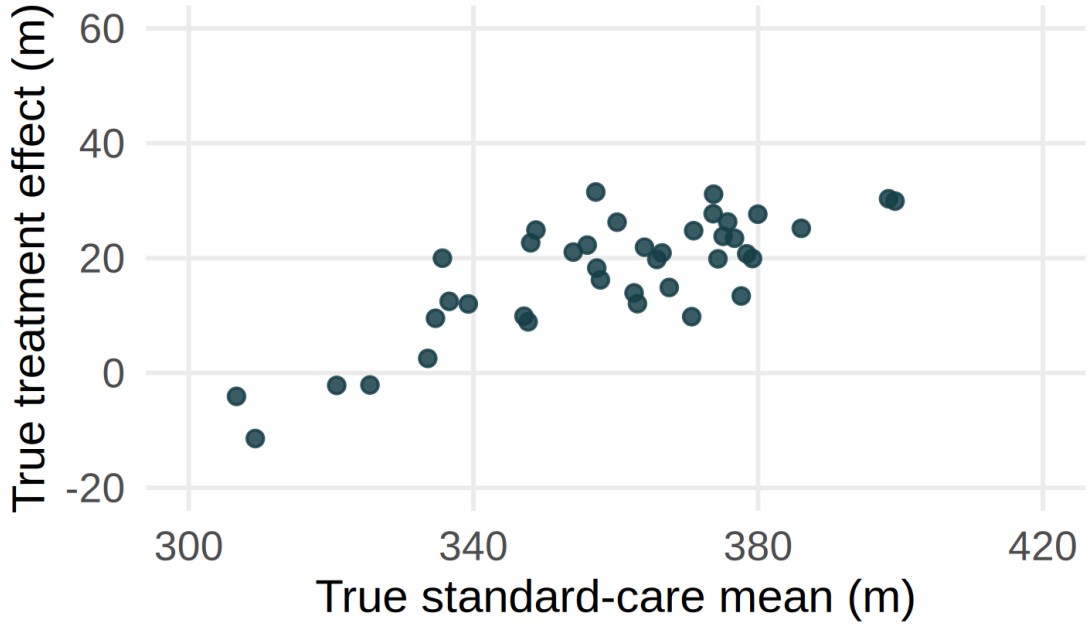
σ_c^2: variance of standard-care means across centers.
σ_d^2: variance of treatment effects across centers.
σ_{cd}: covariance of baseline and treatment-effect deviations.
σ^2: individual variance around the center-specific treatment mean.

Center deviations have population means of zero. They are not constrained to sum to zero in the sampled centers.

How do baseline outcomes and treatment effects vary together?

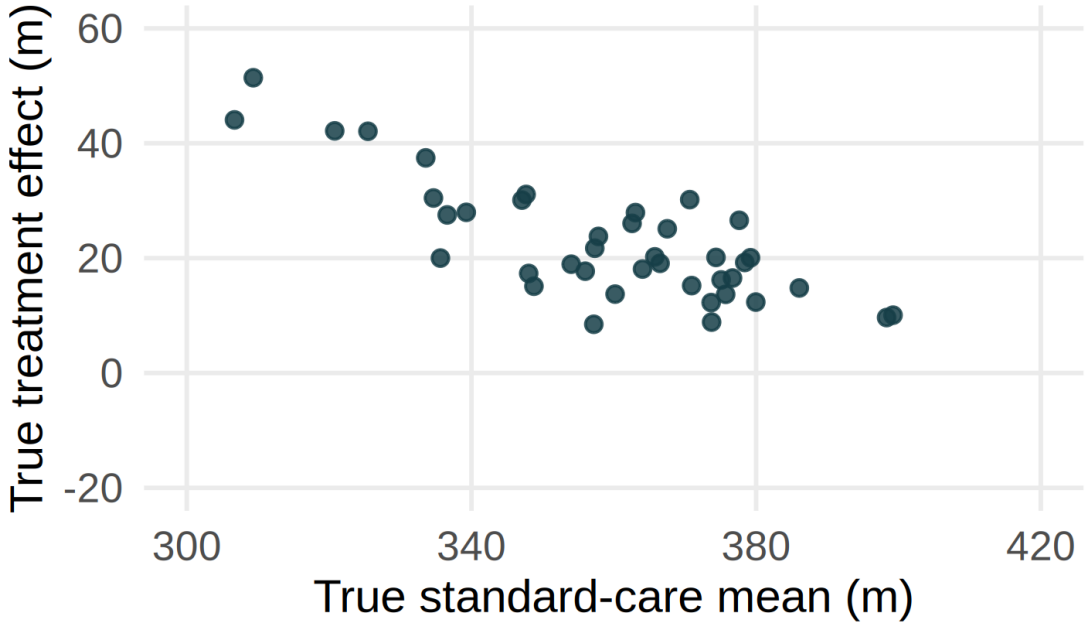
Positive covariance

Higher baseline means tend to accompany larger treatment effects.



Negative covariance

Higher baseline means tend to accompany smaller treatment effects.



True means and enhanced-minus-standard effects, not estimates. Each point is one hypothetical center.

Covariance describes how a center's baseline outcome and treatment effect vary together.

Rehabilitation study: estimated variance components

Results from the random-intercept and random-slope model fitted to the rehabilitation data.

Component	Parameter	Estimate (m²)
Standard-care means across centers	σ_c^2	86.4
Treatment effects across centers	σ_d^2	651.6
Baseline / treatment-effect covariance	σ_{cd}	127.1
Individual residual variation	σ^2	658.4

Estimated baseline / treatment-effect correlation

$$\widehat{\rho}_{cd} = \frac{\widehat{\sigma}_{cd}}{\sqrt{\widehat{\sigma}_c^2 \widehat{\sigma}_d^2}} = \frac{127.1}{\sqrt{86.4 \times 651.6}} \approx 0.54$$

Centers with higher standard-care means tend to have larger treatment effects.

Treatment-effect variation across centers is much larger than baseline variation and almost as large as individual residual variation.

The average treatment effect across the population of centers

The same 120 patients in six centers, fitted with random intercepts and treatment slopes.

Population parameter	Estimate	Estimated SE	95% CI	p-value
Average standard-care outcome (β_0)	349.9	5.0	336.9 to 362.8	< .001
Average treatment effect (β_1)	21.3	11.4	-8.1 to 50.6	0.122

Estimates, estimated SEs and confidence intervals are in meters.

The treatment coefficient estimates the average enhanced-minus-standard effect across the population of centers.

Estimated average benefit: 21.3 m. The 95% confidence interval includes zero.

Same estimated treatment effect, different uncertainty

The balanced design gives the same estimated average treatment effect in both models.

Model	Estimated treatment effect (m)	Estimated SE (m)	p-value
Fixed interaction	21.3	4.68	< .001
Random intercept and slope	21.3	11.43	0.122

Fixed interaction: uncertainty about the average in these six centers comes from sampling patients.

Random intercept and slope: uncertainty about the population average also reflects which centers entered the study.

Treatment effects vary substantially across centers. With only six centers, generalizing to the population average adds considerable uncertainty.

Time for Lab 4



What if we omit the random treatment slope?

Back to rehabilitation: both models target the population-average treatment effect.

Model	Estimated treatment effect (m)	Estimated SE (m)	p-value
Random intercept only	21.3	5.17	< .001
Random intercept and slope	21.3	11.43	0.122

Random intercept only: assumes the same treatment effect in every center.

Random intercept and slope: allows treatment effects to vary across centers.

When treatment effects vary across centers, a random-intercept-only model can underestimate uncertainty about the population-average effect.

Variance decomposition depends on the random-effects structure

WEEK 1: RANDOM INTERCEPT

The same center contribution applies under both treatments.

Outcome variance separates into a center component and an individual component.

This gives one common within-center correlation: the ICC.

WEEK 2: RANDOM INTERCEPT AND SLOPE

The center contribution depends on treatment.

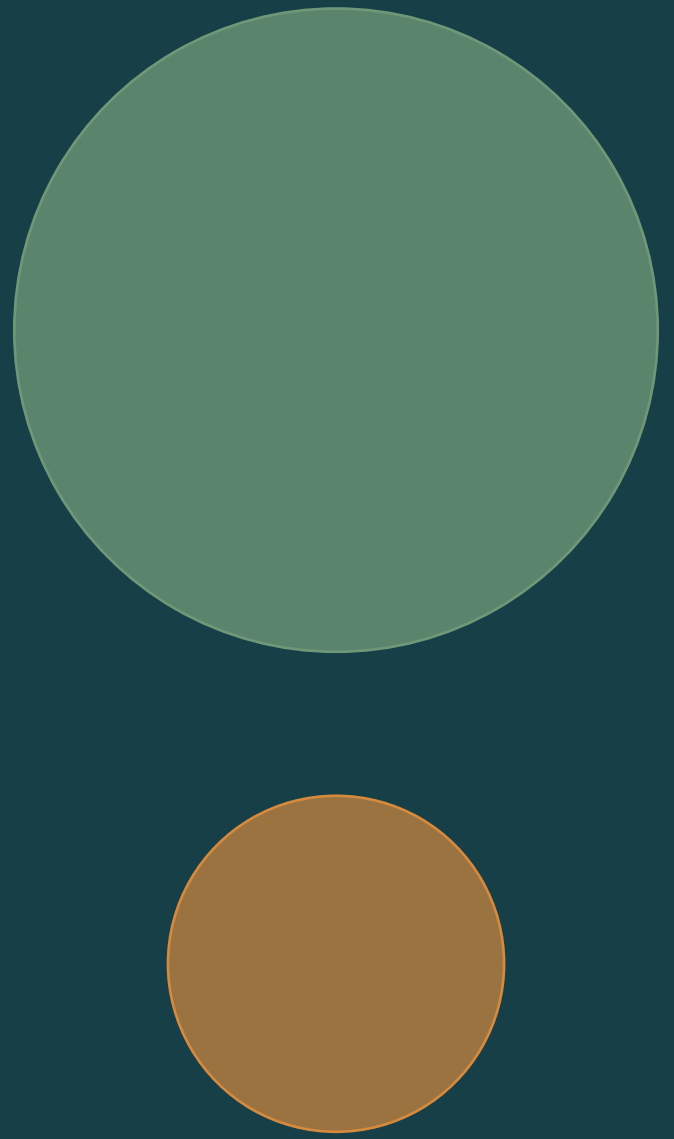
Outcome variance depends on treatment and combines the relevant variances and covariance.

Within-center correlation depends on the treatments received by the two patients.

Variance decomposition still applies, but there is no longer one common variance breakdown or a single ICC for all patients.

PART 3

Crossed random effects



What if there is more than one grouping variable?

So far: one grouping variable, center. What changes with two or more grouping variables?

NESTED GROUPING VARIABLES

Each group at one level belongs to exactly one group at the next level.

Patients belong to wards, and each ward belongs to one hospital.

Grouping variables: ward and hospital.

CROSSED GROUPING VARIABLES

Neither grouping variable is contained within the other.

Each patient is assessed by several radiologists, and each radiologist assesses several patients.

Grouping variables: patient and radiologist.

Next: crossed grouping variables, illustrated by patients and radiologists.
Week 3: nested grouping variables, such as wards within hospitals.

Case study: agreement between radiologists

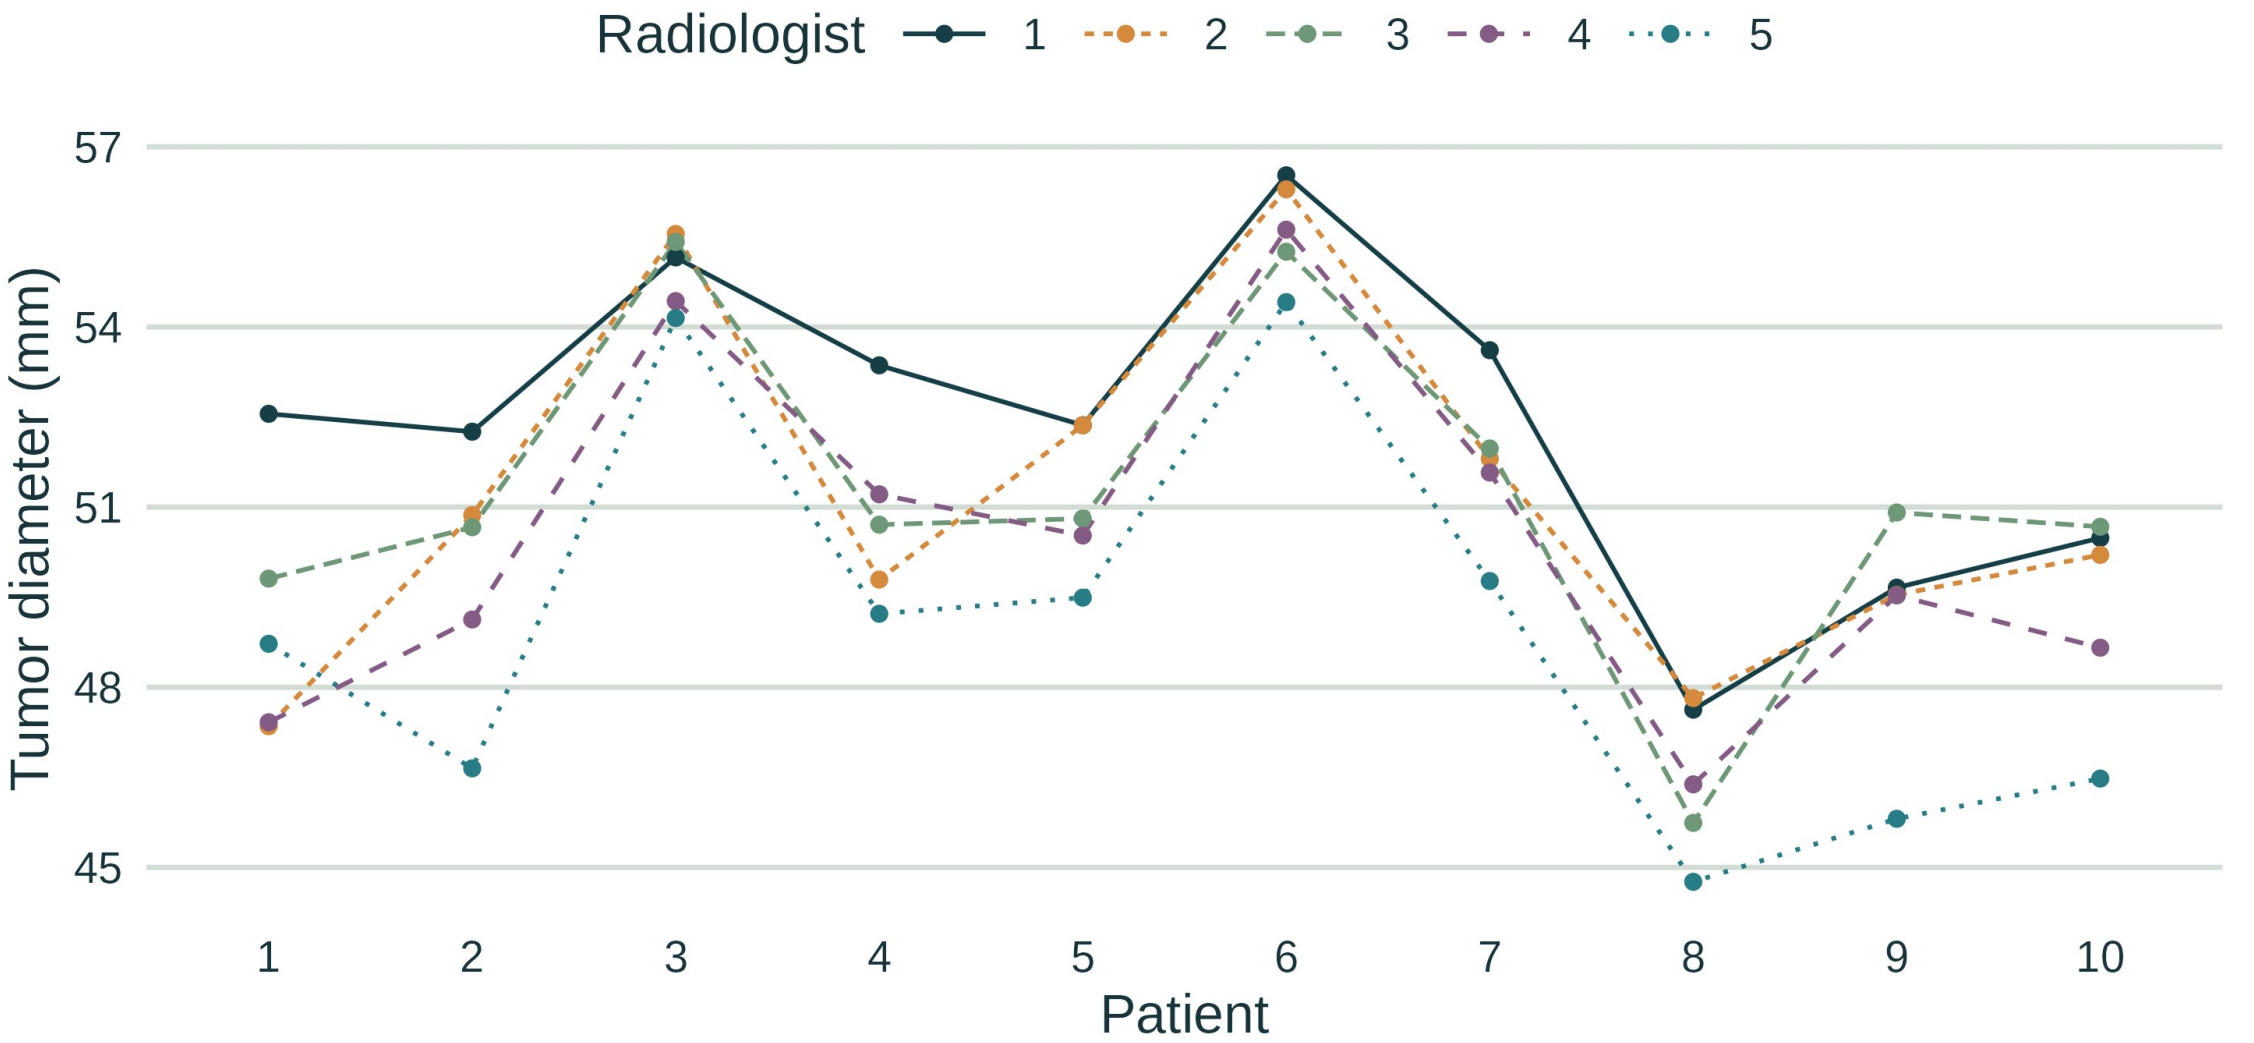
MEASURING TUMOR DIAMETER

- Simulated study: CT scans from 10 patients assessed independently by five radiologists.
- Every radiologist measures the maximum tumor diameter on every scan.
- Outcome: tumor diameter in millimeters, giving 50 measurements.
- We want to generalize beyond these patients and radiologists.

Do different radiologists obtain similar measurements for the same patients?



Exploring the radiologists' measurements



A model with two crossed grouping variables

$$Y_{ij} = \mu + u_i + v_j + \varepsilon_{ij}$$

Y_{ij} is the measurement for patient i made by radiologist j.

WHAT CONTRIBUTES TO A MEASUREMENT?

μ: overall mean tumor diameter.

u_i: patient i's deviation from the mean.

v_j: radiologist j's tendency to measure higher or lower.

ε_{ij}: remaining measurement variation.

TWO RANDOM INTERCEPTS

$$u_i \sim N(0, \sigma^2_{\text{patient}})$$

$$v_j \sim N(0, \sigma^2_{\text{radiologist}})$$

$$\varepsilon_{ij} \sim N(0, \sigma^2)$$

Patient effects, radiologist effects and residual errors are mutually independent.

Measurements are related when they share the same patient or the same radiologist.

Where does measurement variation come from?

Estimated variance components from the crossed model fitted to the radiologist study.

Source of variation	Estimate (mm^2)	Interpretation
Patient	7.44	Differences between patients
Radiologist	1.47	Systematic radiologist differences
Residual	1.00	Remaining measurement variation

$$\widehat{\mathrm{ICC}} = \frac{7.44}{7.44 + 1.47 + 1.00} \approx 0.75$$

Intraclass correlation coefficient (ICC): absolute agreement for a single measurement.

About 75% of the estimated total variation is between patients. Radiologist differences and residual variation reduce agreement.

The End

Questions



